

Filters & Convolution

Computer Vision

Module Code: 600100

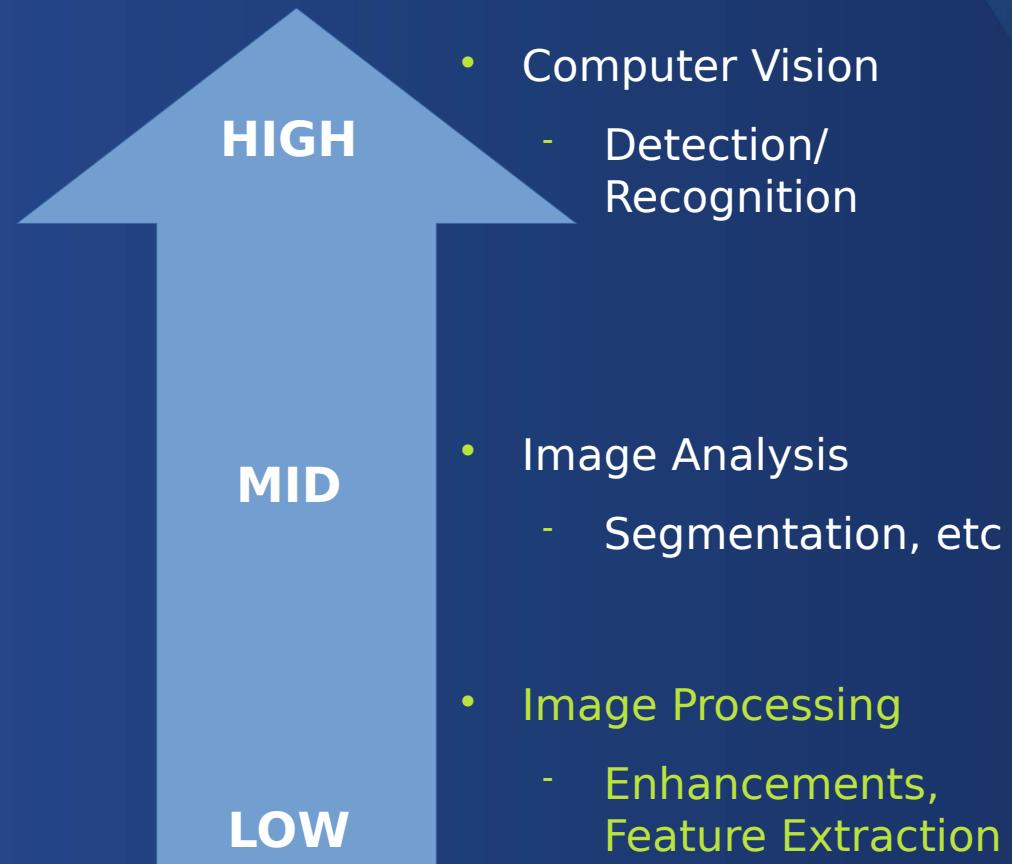
Department of Computer Science and Technology

Today

- Limitations of Point Operations
- From Point to Spatial
- Convolution Operation
 - Have you heard of Convolutional Neural Networks?
 - Padding and Stride



Overview of Computer Vision



- **High Level** - Collection of information in; decision processes
- **Mid Level** – Enhanced images in; Extracted information out.
- **Low Level** – Processing of input image into an ‘enhanced’ state.

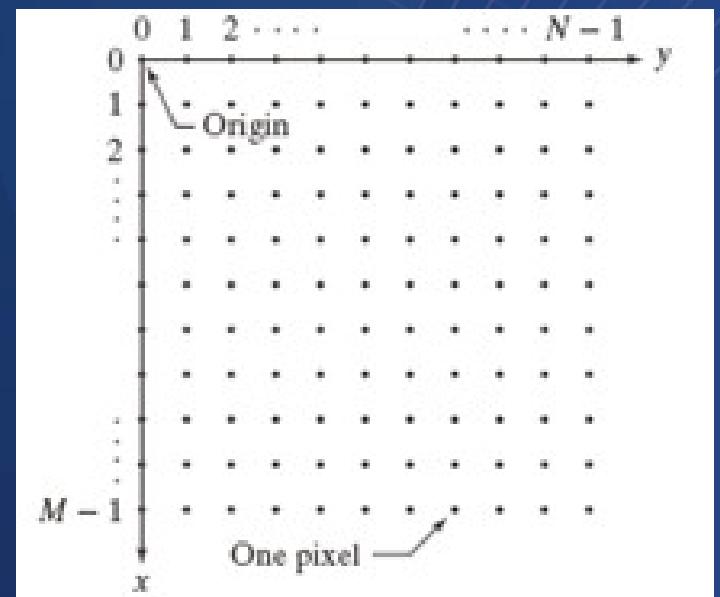
Recap: Image Representation

- Image as a function
- Remember! we sample a real-world object, based on our sensor arrangement.

$$f(x, y) \quad x \in \{0, 1, 2, \dots, M - 1\} \quad \leftarrow \{\text{SET}\}$$
$$y \in \{0, 1, 2, \dots, N - 1\}$$

MEMBER OF

- Intensity value at coords (x,y)
- A Matrix of pixels / pixel elements



Recap: Point Operations

- We can apply arithmetic operations on our image function

$$s(x, y) = f(x, y) + g(x, y)$$

Operations are performed
'Array-like' in a pixel-wise manner

$$d(x, y) = f(x, y) - g(x, y)$$

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} + \begin{bmatrix} 5 & 5 \\ 5 & 5 \end{bmatrix} = \begin{bmatrix} 6 & 7 \\ 8 & 9 \end{bmatrix}$$

$$p(x, y) = f(x, y) \times g(x, y)$$

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 5 & 5 \\ 5 & 5 \end{bmatrix} = \begin{bmatrix} 5 & 10 \\ 15 & 20 \end{bmatrix}$$

$$v(x, y) = f(x, y) \div g(x, y)$$

For all (x,y) pairs!

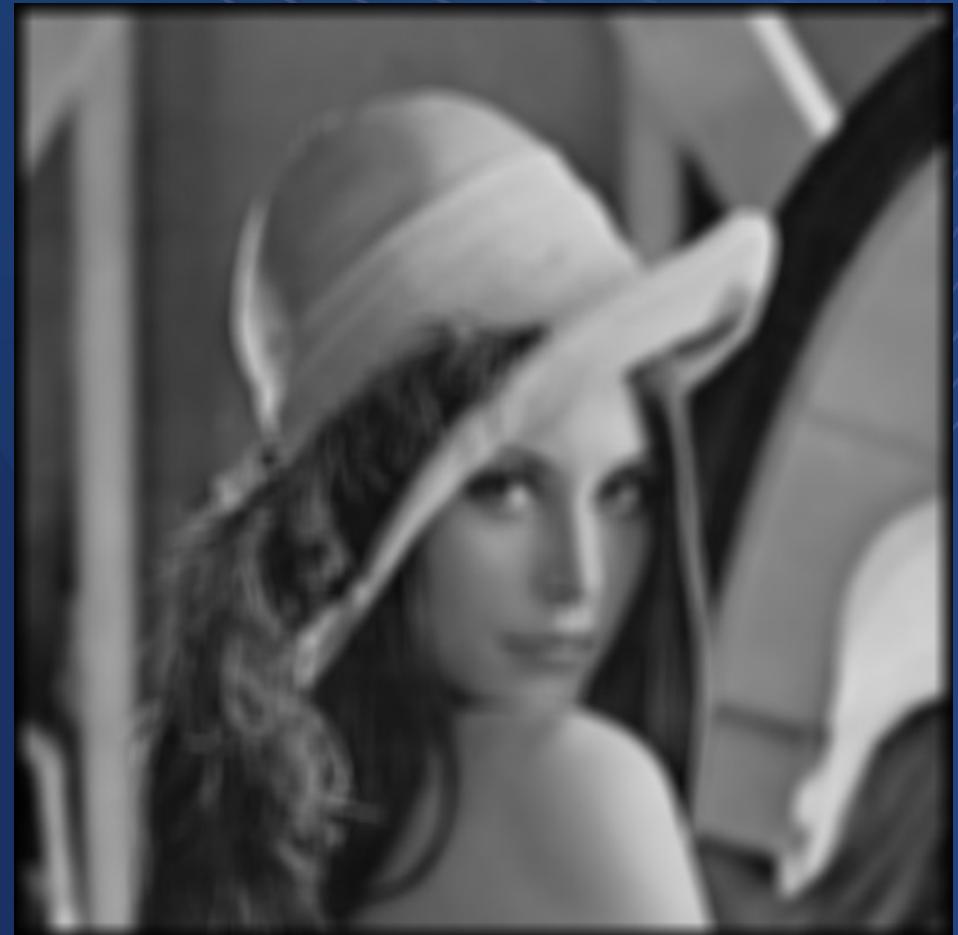
Filtering in the Spatial Domain

Point Operations → Spatial Operations

- Q: What are the limitations of Point Operations?
- A:
 - We can only change each value individually
 - No real context used.
 - We have no information about surrounding points
- Most useful information from images requires more than just a singular value.
 - E.g Edges, Texture, Shapes.
- Therefore, we need to move to the spatial domain to get that information!

Example Limitation

- What can't Point Operations do? Blurring



Example Limitation

- What can't Point Operations do? Sharpening



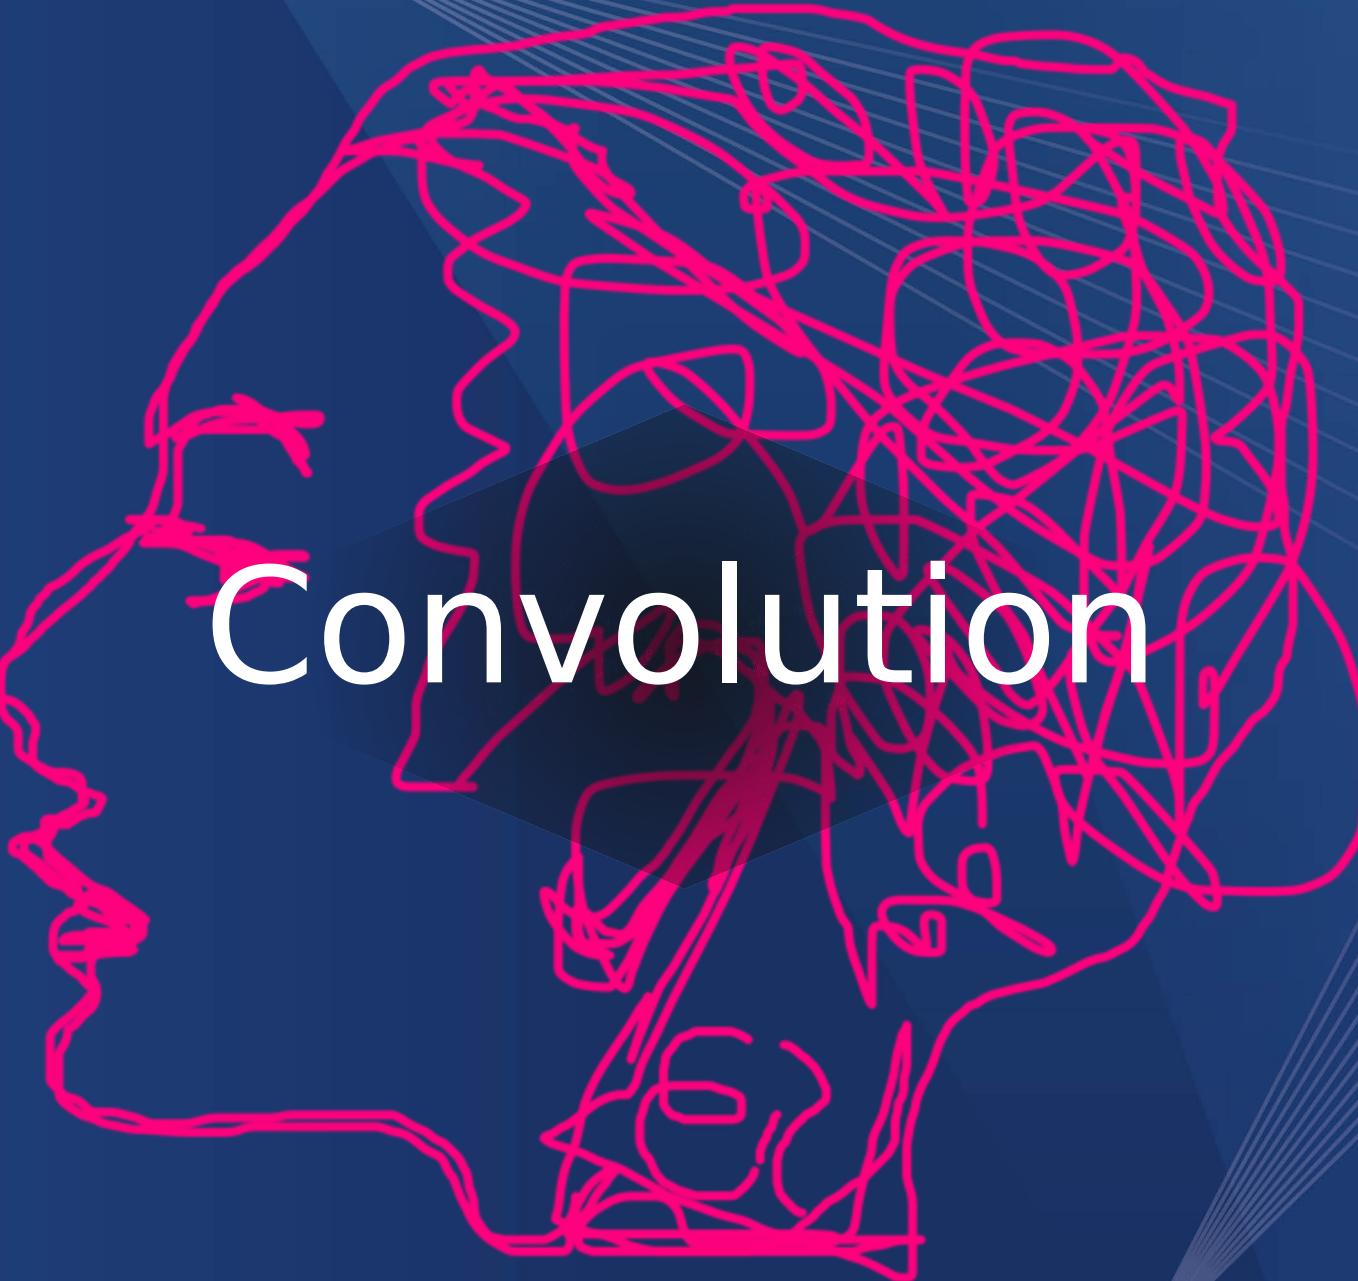
Point Operations → Spatial Operations

- In a similar way to how we can use histograms to provide context, the pixel neighbourhood of a point can tell us a lot of information.
- Q: How could context around a pixel be useful to us, as Computer Vision Detectives?
- Example:
 - Are we a low value in this neighbourhood?
 - How does this intensity value change?
 - Not a lot (low variance), a lot (high variance), not at all?!?
 - Is this region of the image ‘smooth’?

Spatial Filters

- To understand the neighbourhood, we need filters to tell us whether something is ‘important’.
 - A filter is something which passes/modifies/rejects certain inputs (based on their value)
 - Borrowed from the signal processing domain
 - Consists of a separately defined matrix/array from the input image.
- A Spatial Filter makes a decision on a single pixel, based on the neighbourhood.





Convolution

Convolution

- Related to Signal Processing
- Comprised of two components:
 - 1) A Filter (Often called a Kernel or Convolution Mask)
 - 2) Neighbourhood around a Point – x, y.
- Defined as the “Integral of the product of two functions”
 - I.e The input image, and your filter.

$$g(x, y) = \sum_{i=-1}^{i=1} \sum_{j=-1}^{j=1} w(i, j)f(x + i, y + j)$$

Example: 3x3 neighbourhood

1) The Filter

$$w(i, j) = \begin{bmatrix} w_{\frac{-m}{2}, \frac{-n}{2}} & \cdots & w_{\frac{-m}{2}, \frac{+n}{2}} \\ \vdots & & \vdots \\ w_{\frac{+m}{2}, \frac{-n}{2}} & \cdots & w_{\frac{+m}{2}, \frac{+n}{2}} \end{bmatrix}$$

0 ←

All relative to the center

$$= \begin{bmatrix} w_{-1, -1} & w_{-1, 0} & w_{-1, 1} \\ w_{0, -1} & w_{0, 0} & w_{0, 1} \\ w_{1, -1} & w_{1, 0} & w_{1, 1} \end{bmatrix}$$

What is that maths?

$w(i,j) = 3 \times 3$ kernel in this example.

=> -1 → 0 → 1 for both axes

Our filter coefficient at i,j

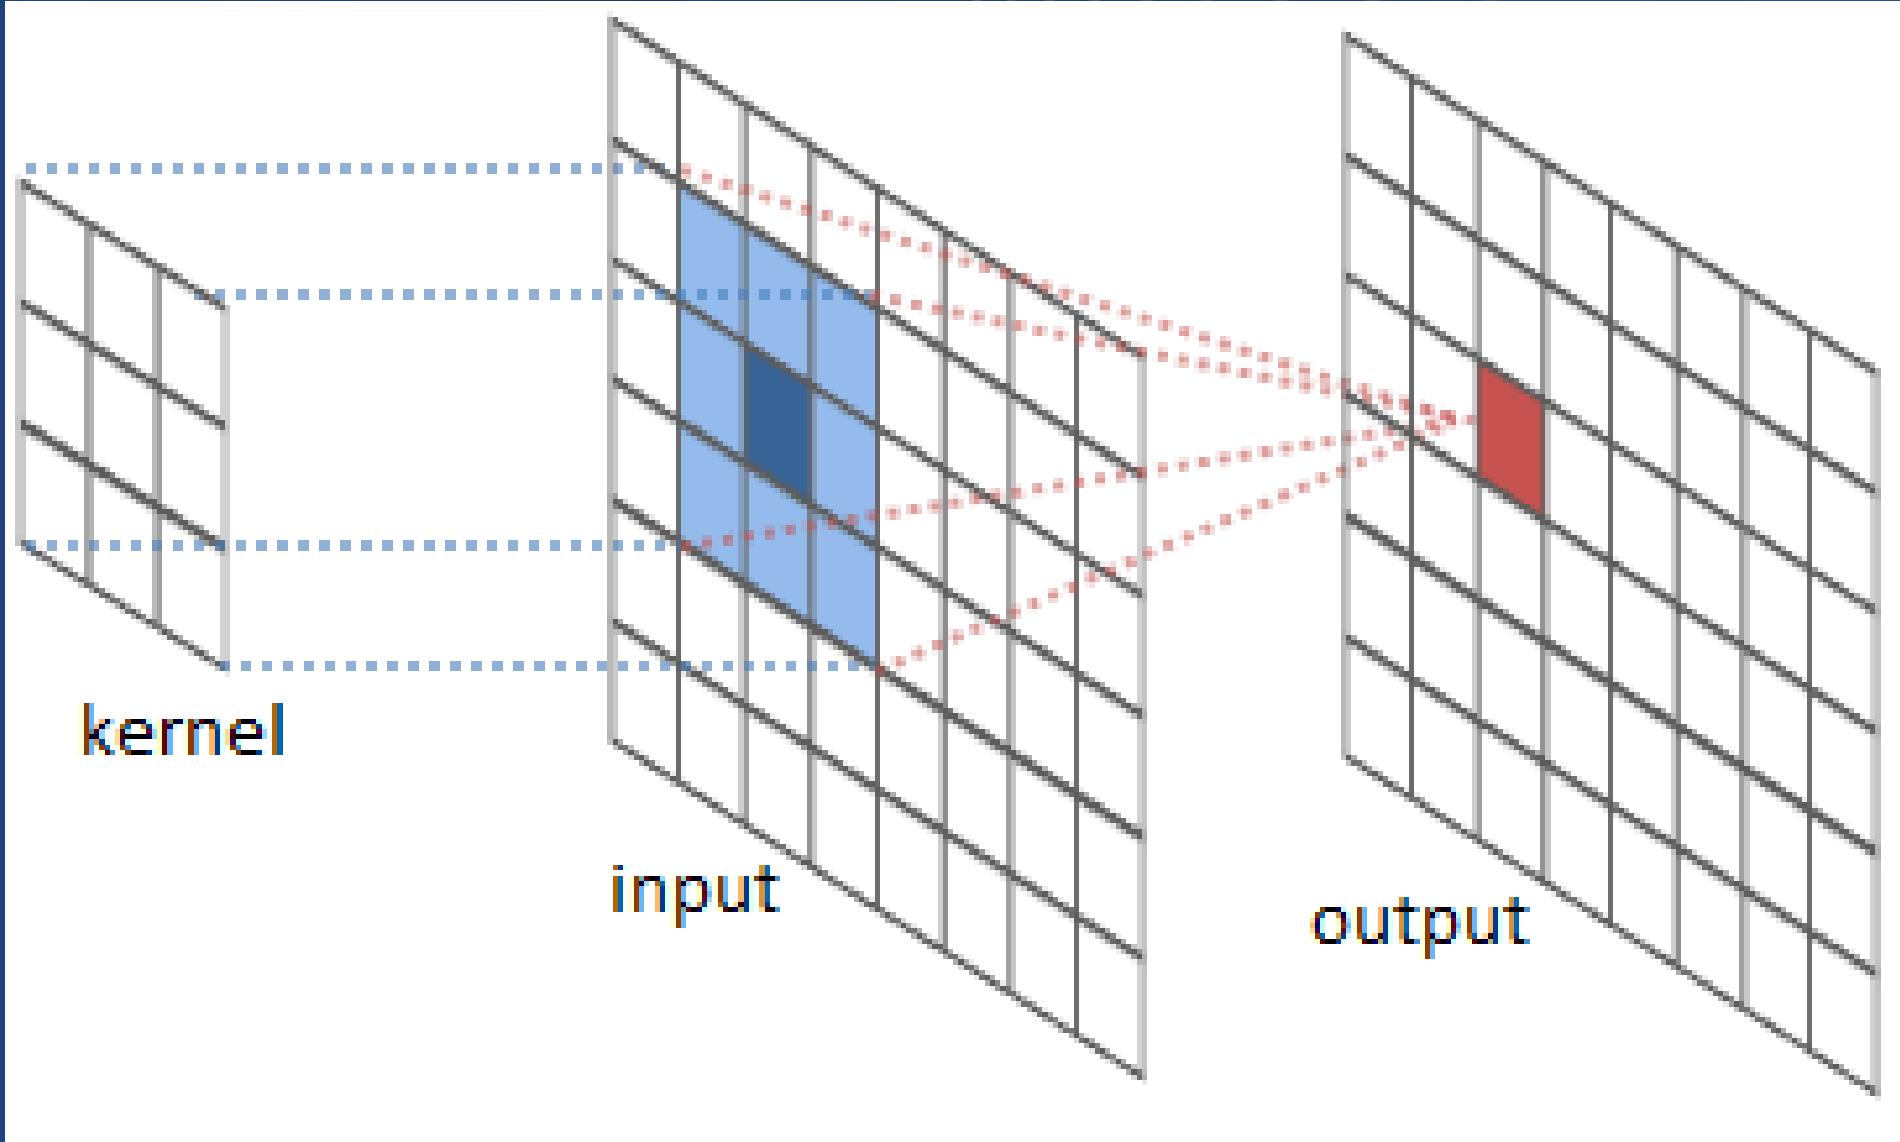
$$g(x, y) = \sum_{i=-1}^{i=1} \sum_{j=-1}^{j=1} w(i, j) f(x + i, y + j)$$

Our 'Convolved' Image

Double Sum
(Double for loop)

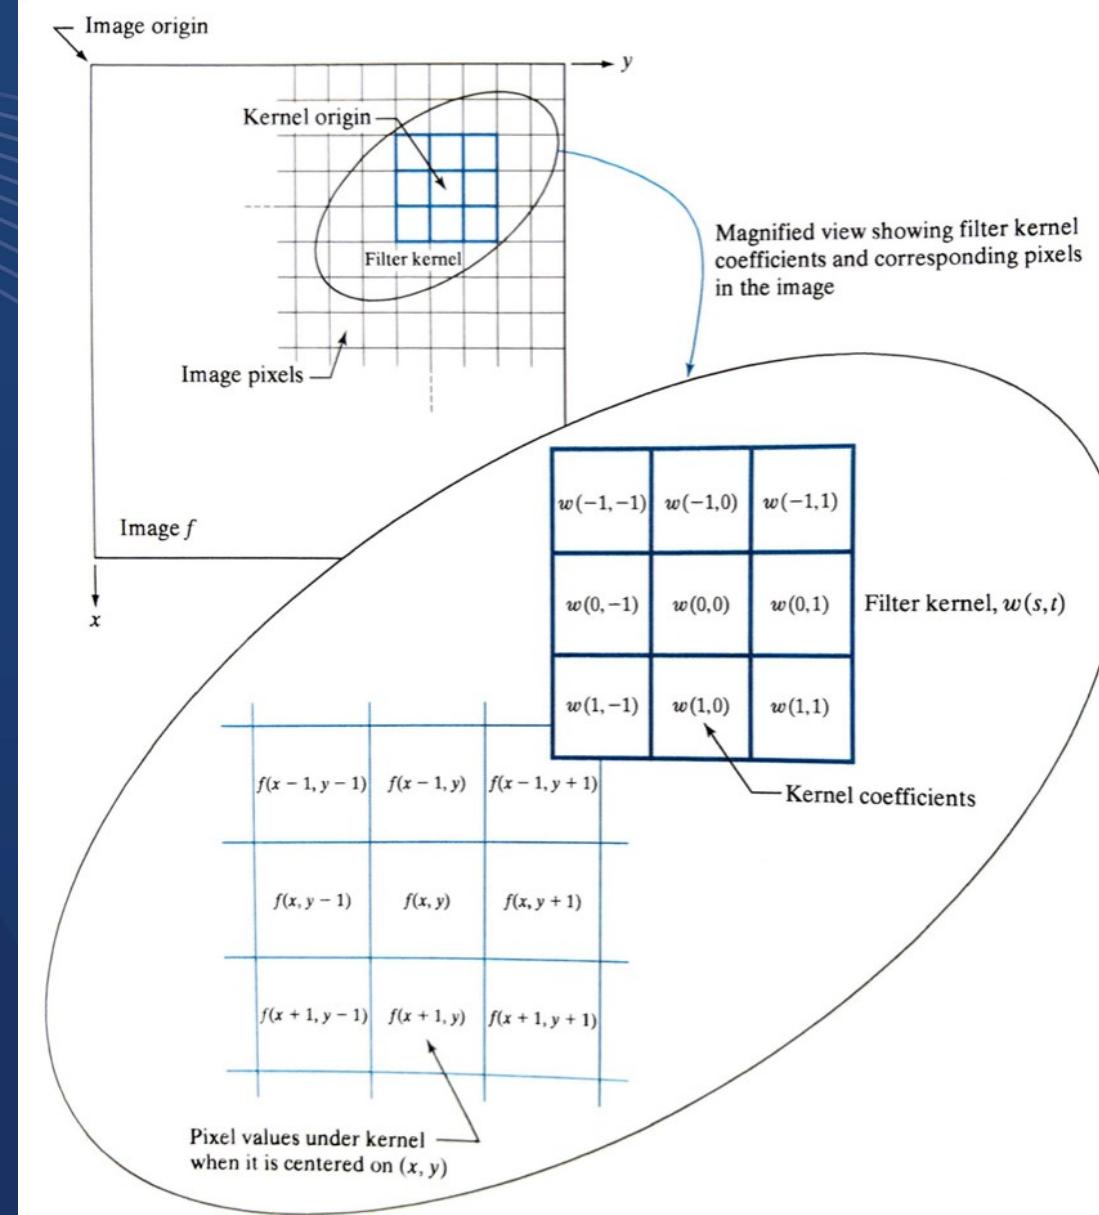
Grey Level $f(x,y)$
Getting neighbourhood via offsets, i, j.

Visual Aid



Today

- Slide the filter (Kernel) over all x, y pairs in the input image
 - For each x, y in $f(x, y)$.
- Convolve the input image neighbourhood, with the filter
 - Filter is centered around a given x, y
- Generates a new output image, $g(x, y)$
 - Considered the weighted sum of the neighbourhood.



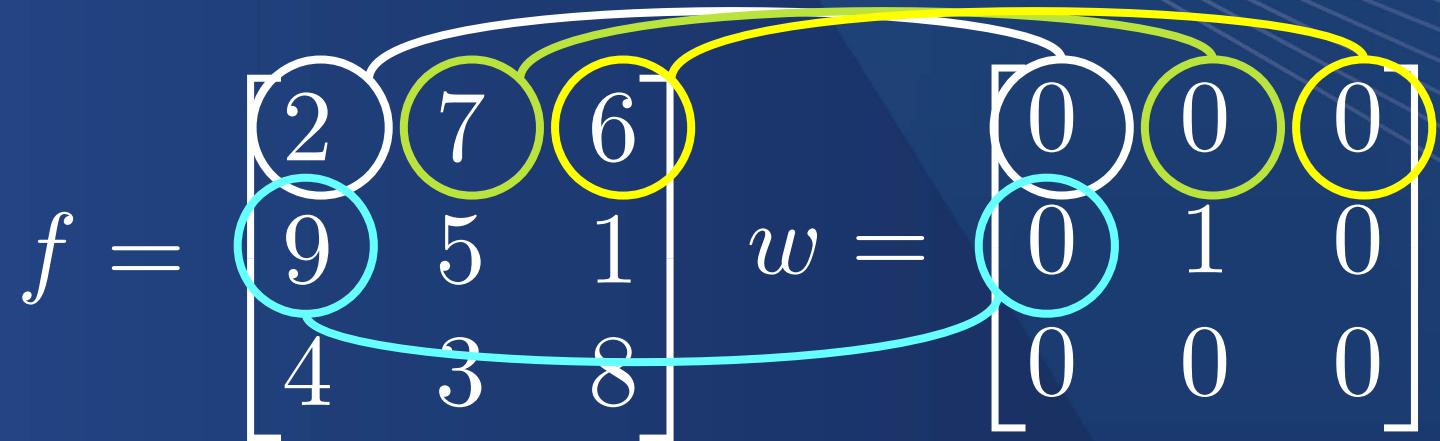
Convolution Example 1

$$f = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix} \quad w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

$$f \star w = ?$$

f convolved with w

Convolution Example 1



$$\begin{aligned} & (0 \times 2) + (0 \times 7) + (0 \times 6) && 0 + 0 + 0 \\ = & + (0 \times 9) + (1 \times 5) + (0 \times 1) &= + 0 + 5 + 0 = 5 \\ & + (0 \times 4) + (0 \times 3) + (0 \times 8) && + 0 + 0 + 0 \end{aligned}$$

Convolution Example 2

$$f = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix} \quad w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

$$f \star w = ?$$

Convolution Example 2

$$f = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix} \quad w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

$$\begin{aligned} & (0 \times 2) + (0 \times 7) + (0 \times 6) && 0 + 0 + 0 \\ = & + (0 \times 9) + (0 \times 5) + (1 \times 1) = + 0 + 0 + 1 = 1 \\ & + (0 \times 4) + (0 \times 3) + (0 \times 8) && + 0 + 0 + 0 \end{aligned}$$

Convolution Example 3

$$f = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix} \quad w = \begin{bmatrix} \frac{1}{4} & 0 & \frac{1}{4} \\ 0 & 0 & 0 \\ \frac{1}{4} & 0 & \frac{1}{4} \end{bmatrix}$$

$$f \star w = ?$$

Convolution Example 3

$$f = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix} \quad w = \begin{bmatrix} \frac{1}{4} & 0 & \frac{1}{4} \\ 0 & 0 & 0 \\ \frac{1}{4} & 0 & \frac{1}{4} \end{bmatrix}$$

$$\begin{aligned} & (\frac{1}{4} \times 2) + (0 \times 7) + (\frac{1}{4} \times 6) & \frac{1}{2} + 0 + \frac{3}{2} & +2 \\ & + (0 \times 9) + (0 \times 5) + (0 \times 1) & + 0 + 0 + 0 & +0 \\ & + (\frac{1}{4} \times 4) + (0 \times 3) + (\frac{1}{4} \times 8) & + 1 + 0 + 2 & +3 \\ & & & \hline & & & 5 \end{aligned}$$

Boundary Conditions

- Convolving 3x3 neighbourhood with a 3x3 filter
 - Produces a 1x1 output.
- How can we apply this to a whole image?
 - Sliding the filter through the original input image.
 - Filter does ‘fit’ perfectly into the very corner pixels
 - Filter would ‘hang off’ the page...
- Solution: Padding!



Padding



No Padding

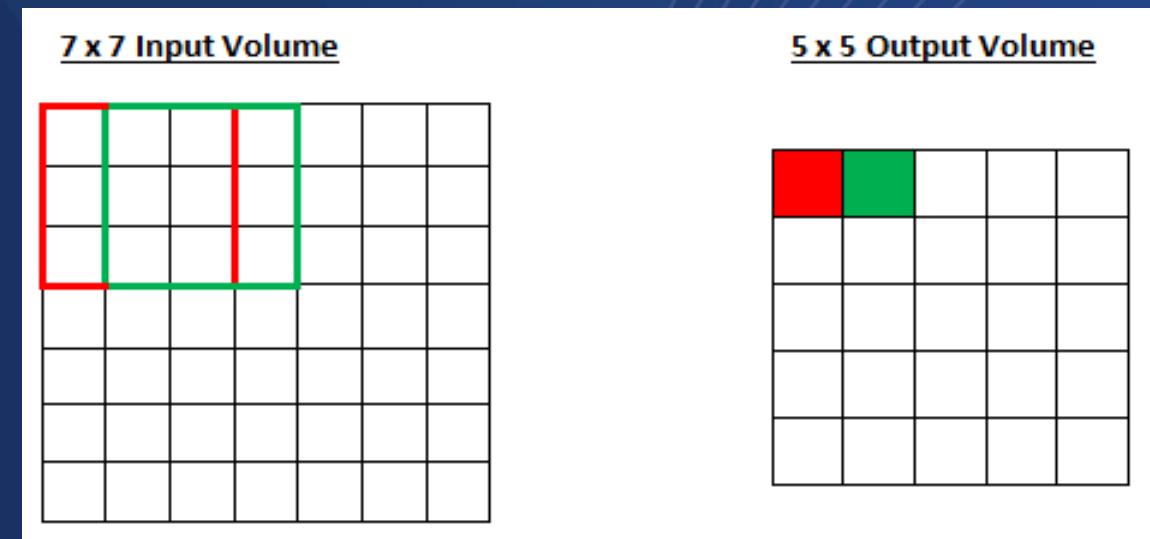
- Simplest form of padding is No Padding
- If we can't center the filter on all our input x,y pairs, that's okay.
 - → We will put the filter as close as possible for boundaries.
- Results in an overall decrease in image size
 - E.g $7 \times 7 \rightarrow 5 \times 5$ after Convolution.

We lose ceil(Filter Width / 2)

E.g 3×3 filter $\rightarrow \text{ceil}(3/2) = \text{ceil}(1.5) = 2$

Total loss of 2.

Per-axis, as it is a square matrix.

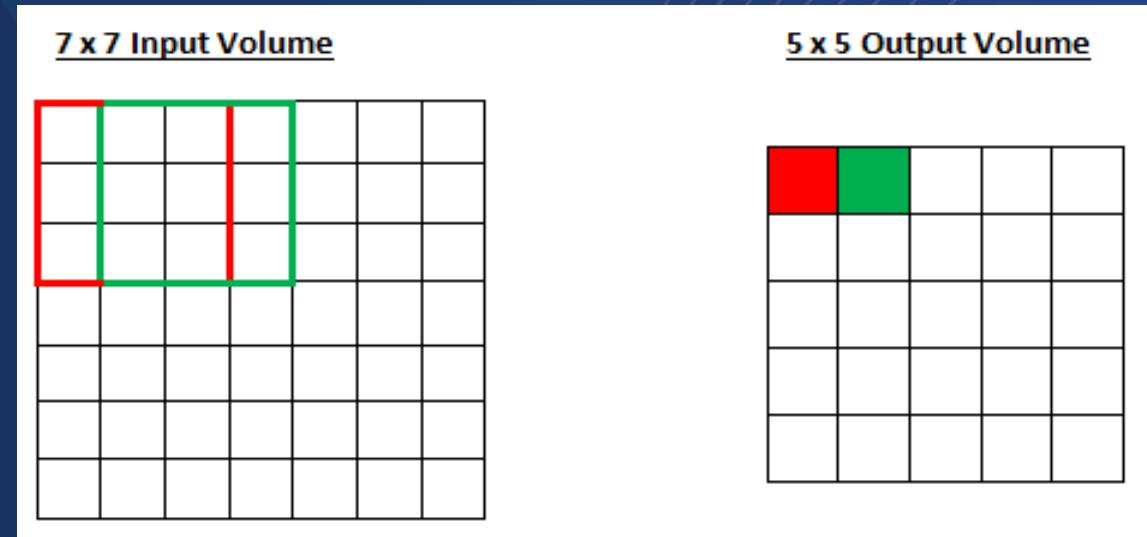


'Same' Padding (sometimes called full / zero)

- Provides the 'same' output resolution, as input. (full padding in MATLAB)
- Pads the original image with 0 around all sides.
 - Amount of zeros depends on how big our kernel is.
- Typically denoted by:

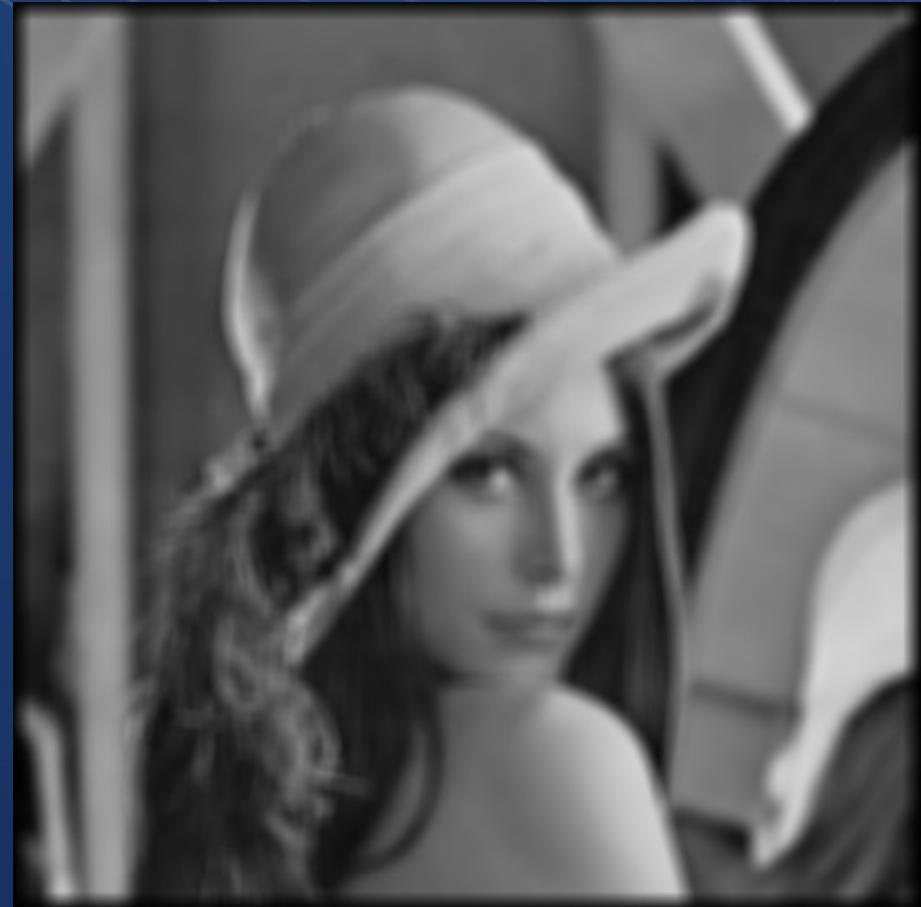
$$Padding = \frac{FilterSize - 1}{2}$$

- Padding = $(3 - 1) / 2 = 1$ pixel padding
- This is intuitive.



Padding Artifacts

- Padding edges with zeros adds a dark ‘border’
- Depending on kernel size, this may have an impact on the output result.
- If kernel weights are non-central:
 - Lots of those zeros will be used in the full convolution calculation
 - → Massively skewed values.
- E.g Zero Padding causes Black Borders to appear in our ‘blur’ example.



Zero Padding

$$f = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 7 & 6 & 0 \\ 0 & 9 & 5 & 1 & 0 \\ 0 & 4 & 3 & 8 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

$$f \star w = \begin{bmatrix} 2 & 7 & 6 \\ 9 & 5 & 1 \\ 4 & 3 & 8 \end{bmatrix}$$

Stride

- Filters so far have moved one pixel at a time
 - E.g Stride = 1.
- Stride represents how far to move the filter each time.
- Higher Stride → Smaller Output (Think sampling)
- Q: Why do this?
- A: Quicker, Cheaper, Less Memory Usage
- A: We might not need to do it on every pixel, depends on our pixel resolution, and how detailed our subject is.



Zero Padding w/ Stride = 2

$$f = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 7 & 6 & 0 \\ 0 & 9 & 5 & 1 & 0 \\ 0 & 4 & 3 & 8 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

The input feature map f is a 5x5 matrix with values: (0,0) = 0, (0,1) = 0, (0,2) = 0, (0,3) = 0, (0,4) = 0, (1,0) = 0, (1,1) = 2, (1,2) = 7, (1,3) = 6, (1,4) = 0, (2,0) = 0, (2,1) = 9, (2,2) = 5, (2,3) = 1, (2,4) = 0, (3,0) = 0, (3,1) = 4, (3,2) = 3, (3,3) = 8, (3,4) = 0, (4,0) = 0, (4,1) = 0, (4,2) = 0, (4,3) = 0, (4,4) = 0. It is zero-padded with a stride of 2.

$$w = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

The weight w is a 3x3 matrix with values: (0,0) = 0, (0,1) = 0, (0,2) = 0, (1,0) = 0, (1,1) = 1, (1,2) = 0, (2,0) = 0, (2,1) = 0, (2,2) = 0.

$$f \star w =$$


The result of the convolution $f \star w$ is a 2x2 matrix with values: (0,0) = 2, (0,1) = 6, (1,0) = 4, (1,1) = 8. The values are colored green for the top-left element (2), red for the bottom-left element (4), green for the top-right element (6), and red for the bottom-right element (8).



Questions? (now, after, or e-mail)

